

Task-1

Accessing the Data

```
In [1]: #KPMG provides us multiple Dataset, So loading them first
import pandas as pd
import numpy as np
data = pd.ExcelFile("C:/Users/91913/Downloads/KPMG_VI_New_raw_data_update_final (2).xlsx")
```

Reading each file

```
In [2]: Transactions=pd.read_excel(data,'Transactions')
CustomerDemographic = pd.read_excel(data, 'CustomerDemographic')
CustomerAddress = pd.read_excel(data, 'CustomerAddress')
```

Deleting the names and setting row 1 as the columns

```
In [3]: Transactions.head()
```

Out[3]:

Note: The data and information in this document is reflective of a hypothetical situation and client. This document is to be used for KPMG Virtual Internship purposes only.

	Unnamed: 1	Unnamed: 2	Unnamed: 3	Unnamed: 4	Unnamed: 5	Unnamed: 6	Unnamed: 7	Unna	
0	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	product_line	produ
1	1	2	2950	2017-02-25 00:00:00	False	Approved	Solex	Standard	r
2	2	3	3120	2017-05-21 00:00:00	True	Approved	Trek Bicycles	Standard	r
3	3	37	402	2017-10-16 00:00:00	False	Approved	OHM Cycles	Standard	
4	4	88	3135	2017-08-31 00:00:00	False	Approved	Norco Bicycles	Standard	r

```
In [4]: Transactions.columns=Transactions.iloc[0]
```

```
In [5]: Transactions.drop(0,axis=0,inplace=True)
```

First look

```
In [6]: Transactions
```

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	product_line
1	1	2	2950	2017-02-25 00:00:00	False	Approved	Solex	Standard
2	2	3	3120	2017-05-21 00:00:00	True	Approved	Trek Bicycles	Standard
3	3	37	402	2017-10-16 00:00:00	False	Approved	OHM Cycles	Standard
4	4	88	3135	2017-08-31 00:00:00	False	Approved	Norco Bicycles	Standard
5	5	78	787	2017-10-01 00:00:00	True	Approved	Giant Bicycles	Standard
...
19996	19996	51	1018	2017-06-24 00:00:00	True	Approved	OHM Cycles	Standard
19997	19997	41	127	2017-11-09 00:00:00	True	Approved	Solex	Road
19998	19998	87	2284	2017-04-14 00:00:00	True	Approved	OHM Cycles	Standard
19999	19999	6	2764	2017-07-03 00:00:00	False	Approved	OHM Cycles	Standard
20000	20000	11	1144	2017-09-22 00:00:00	True	Approved	Trek Bicycles	Standard

20000 rows × 13 columns

Looking at the dimensionality

```
In [7]: Transactions.shape
```

```
Out[7]: (20000, 13)
```

Looking at the info

```
In [8]: Transactions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 1 to 20000
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   transaction_id      20000 non-null  object
1   product_id          20000 non-null  object
2   customer_id         20000 non-null  object
3   transaction_date     20000 non-null  object
```

```

4   online_order      19640 non-null object
5   order_status      20000 non-null object
6   brand             19803 non-null object
7   product_line      19803 non-null object
8   product_class     19803 non-null object
9   product_size      19803 non-null object
10  list_price        20000 non-null object
11  standard_cost      19803 non-null object
12  product_first_sold_date 19803 non-null object
dtypes: object(13)
memory usage: 2.0+ MB

```

By looking at the info, the dtype of many columns are not correct. We will change them further.

Looking at the missing values

```
In [9]: Transactions.isna().sum()
```

```

Out[9]: 0
transaction_id      0
product_id          0
customer_id         0
transaction_date    0
online_order       360
order_status        0
brand              197
product_line        197
product_class       197
product_size        197
list_price          0
standard_cost       197
product_first_sold_date 197
dtype: int64

```

```

In [10]: Transactions["online_order"].fillna("na",inplace=True)
Transactions["brand"].fillna("not known",inplace=True)
Transactions["product_line"].fillna("not known",inplace=True)
Transactions["product_class"].fillna("not known",inplace=True)
Transactions["product_size"].fillna("not known",inplace=True)
Transactions["standard_cost"].fillna(Transactions["standard_cost"].mean(),inplace=True)
Transactions["product_first_sold_date"].fillna("0",inplace=True)

```

```
In [11]: Transactions.isnull().sum()
```

```

Out[11]: 0
transaction_id      0
product_id          0
customer_id         0
transaction_date    0
online_order        0
order_status        0
brand              0
product_line        0
product_class       0
product_size        0
list_price          0
standard_cost       0
product_first_sold_date 0
dtype: int64

```

```
In [12]: Transactions.duplicated().sum()
```

```

Out[12]: 0

```

There are no duplicate values, all the data is unique.

Recasting the Data type

```
In [13]: Transactions=Transactions.astype({"transaction_id":"int64",
                                           "product_id":"int64",
                                           "customer_id":"int64",
                                           "list_price":"int64",
                                           "standard_cost":"int64",

                                           })
```

```
In [14]: Transactions['transaction_date'] = pd.to_datetime(Transactions['transaction_date'])
```

```
In [15]: Transactions
```

```
Out[15]:
```

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	product_line
1	1	2	2950	2017-02-25	False	Approved	Solex	Standard
2	2	3	3120	2017-05-21	True	Approved	Trek Bicycles	Standard
3	3	37	402	2017-10-16	False	Approved	OHM Cycles	Standard
4	4	88	3135	2017-08-31	False	Approved	Norco Bicycles	Standard
5	5	78	787	2017-10-01	True	Approved	Giant Bicycles	Standard
...
19996	19996	51	1018	2017-06-24	True	Approved	OHM Cycles	Standard
19997	19997	41	127	2017-11-09	True	Approved	Solex	Road
19998	19998	87	2284	2017-04-14	True	Approved	OHM Cycles	Standard
19999	19999	6	2764	2017-07-03	False	Approved	OHM Cycles	Standard
20000	20000	11	1144	2017-09-22	True	Approved	Trek Bicycles	Standard

20000 rows × 13 columns

```
In [16]: Transactions.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 1 to 20000
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
```

```

0    transaction_id      20000 non-null int64
1    product_id         20000 non-null int64
2    customer_id        20000 non-null int64
3    transaction_date    20000 non-null datetime64[ns]
4    online_order        20000 non-null object
5    order_status        20000 non-null object
6    brand               20000 non-null object
7    product_line        20000 non-null object
8    product_class       20000 non-null object
9    product_size        20000 non-null object
10   list_price          20000 non-null int64
11   standard_cost       20000 non-null int64
12   product_first_sold_date 20000 non-null object
dtypes: datetime64[ns](1), int64(5), object(7)
memory usage: 2.0+ MB

```

looking at the unique values in the Data

```
In [17]: Transactions.nunique()
```

```

Out[17]:
0
transaction_id      20000
product_id         101
customer_id        3494
transaction_date    364
online_order        3
order_status        2
brand              7
product_line        5
product_class       4
product_size        4
list_price         274
standard_cost       97
product_first_sold_date 101
dtype: int64

```

Exploring the columns

```
In [18]: Transactions.columns
```

```

Out[18]:
Index(['transaction_id', 'product_id', 'customer_id', 'transaction_date',
      'online_order', 'order_status', 'brand', 'product_line',
      'product_class', 'product_size', 'list_price', 'standard_cost',
      'product_first_sold_date'],
      dtype='object', name=0)

```

```
In [19]: Transactions['brand'].value_counts()
```

```

Out[19]:
Solex      4253
Giant Bicycles 3312
WeareA2B    3295
OHM Cycles  3043
Trek Bicycles 2990
Norco Bicycles 2910
not known   197
Name: brand, dtype: int64

```

```
In [20]: Transactions['order_status'].value_counts()
```

```

Out[20]:
Approved    19821
Cancelled    179
Name: order_status, dtype: int64

```

```
In [21]: Transactions['product_line'].value_counts()

Out[21]: Standard      14176
Road      3970
Touring      1234
Mountain      423
not known      197
Name: product_line, dtype: int64

In [22]: Transactions['product_class'].value_counts()

Out[22]: medium      13826
high      3013
low      2964
not known      197
Name: product_class, dtype: int64

In [23]: Transactions['online_order'].value_counts()

Out[23]: True      9829
False      9811
na      360
Name: online_order, dtype: int64

In [24]: Transactions['product_first_sold_date'].value_counts()

Out[24]: 33879      234
41064      229
37823      227
39880      222
38216      220
...
41848      169
42404      168
41922      166
37659      163
34586      162
Name: product_first_sold_date, Length: 101, dtype: int64

In [ ]:
```

Adding profit and profit% columns

```
In [25]: Transactions["profit"]=Transactions["list_price"]-Transactions["standard_cost"]
Transactions["profit_percentage"]=(Transactions["list_price"]-Transactions["standard_cos
```

Exploring the Data in Customer Demographic Data set

```
In [26]: cd=pd.read_excel(data,"CustomerDemographic")

In [27]: cd
```

a
hypothetical
situation
and client.
This
document is
to be used
for KPMG
Virtual
Internship
purposes
only.

0	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title
1	1	Laraine	Medendorp	F	93	1953-10-12 00:00:00	Executive Secretary
2	2	Eli	Bockman	Male	81	1980-12-16 00:00:00	Administrative Officer
3	3	Arlin	Dearle	Male	61	1954-01-20 00:00:00	Recruiting Manager
4	4	Talbot	NaN	Male	33	1961-10-03 00:00:00	NaN
...
3996	3996	Rosalia	Halgarth	Female	8	1975-08-09 00:00:00	VP Product Management
3997	3997	Blanch	Nisuis	Female	87	2001-07-13 00:00:00	Statistician I
3998	3998	Sarene	Woolley	U	60	NaN	Assistant Manager
3999	3999	Patrizius	NaN	Male	11	1973-10-24 00:00:00	NaN
4000	4000	Kippy	Oldland	Male	76	1991-11-05 00:00:00	Software Engineer IV

4001 rows × 13 columns

```
In [28]: cd.columns=cd.iloc[0]
```

```
In [29]: cd=cd.drop(0,axis=0)
```

```
In [30]: cd.tail()
```

Out[30]:	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	job
	3996	3996	Rosalia	Halgarth	Female	8	1975-08-09 00:00:00	VP Product Management

3997	3997	Blanch	Nisuis	Female	87	2001-07-13 00:00:00	Statistician II
3998	3998	Sarene	Woolley	U	60	NaN	Assistant Manager
3999	3999	Patrizius	NaN	Male	11	1973-10-24 00:00:00	NaN
4000	4000	Kippy	Oldland	Male	76	1991-11-05 00:00:00	Software Engineer IV

In [31]:

cd.shape

Out[31]:

(4000, 13)

In [32]:

cd.duplicated().sum()

Out[32]:

0

No duplicate value presents

In [33]:

cd.describe().T

Out[33]:

	count	unique	top	freq
0				
customer_id	4000	4000	1	1
first_name	4000	3139	Max	5
last_name	3875	3725	Pristnor	3
gender	4000	6	Female	2037
past_3_years_bike_related_purchases	4000	100	16	56
DOB	3913	3448	1978-01-30 00:00:00	7
job_title	3494	195	Business Systems Development Analyst	45
job_industry_category	3344	9	Manufacturing	799
wealth_segment	4000	3	Mass Customer	2000
deceased_indicator	4000	2	N	3998
default	3698	90	100	113
owns_car	4000	2	Yes	2024
tenure	3913	22	7	235

In [34]:

cd.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4000 entries, 1 to 4000

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
----	-----	-----	-----
0	customer_id	4000 non-null	object
1	first_name	4000 non-null	object

2	last_name	3875	non-null	object
3	gender	4000	non-null	object
4	past_3_years_bike_related_purchases	4000	non-null	object
5	DOB	3913	non-null	object
6	job_title	3494	non-null	object
7	job_industry_category	3344	non-null	object
8	wealth_segment	4000	non-null	object
9	deceased_indicator	4000	non-null	object
10	default	3698	non-null	object
11	owns_car	4000	non-null	object
12	tenure	3913	non-null	object

dtypes: object(13)
memory usage: 406.4+ KB

unique values

```
In [35]: cd.nunique()
```

```
Out[35]: 0
customer_id          4000
first_name           3139
last_name            3725
gender                6
past_3_years_bike_related_purchases  100
DOB                  3448
job_title            195
job_industry_category    9
wealth_segment         3
deceased_indicator       2
default              90
owns_car               2
tenure                22
dtype: int64
```

missing values

```
In [36]: cd.isnull().sum()
```

```
Out[36]: 0
customer_id          0
first_name           0
last_name            125
gender                0
past_3_years_bike_related_purchases  0
DOB                  87
job_title            506
job_industry_category  656
wealth_segment         0
deceased_indicator       0
default              302
owns_car               0
tenure                87
dtype: int64
```

```
In [37]: cd["last_name"].fillna("not_known",inplace=True)
cd["job_title"].fillna("not_known",inplace=True)
cd["job_industry_category"].fillna("not_known",inplace=True)
cd["tenure"].fillna(int(0),inplace=True)
```

```
In [38]: cd.isnull().sum()
```

```
Out[38]: 0
customer_id 0
first_name 0
last_name 0
gender 0
past_3_years_bike_related_purchases 0
DOB 87
job_title 0
job_industry_category 0
wealth_segment 0
deceased_indicator 0
default 302
owns_car 0
tenure 0
dtype: int64
```

We will delete the Column "default" because of having non-readable values.

```
In [39]: cd.drop("default",axis=1,inplace=True)
```

```
In [40]: cd
```

Out[40]:

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	j
1	1	Laraine	Medendorp	F	93	1953-10-12 00:00:00	Executive Secretary	
2	2	Eli	Bockman	Male	81	1980-12-16 00:00:00	Administrative Officer	
3	3	Arlin	Dearle	Male	61	1954-01-20 00:00:00	Recruiting Manager	
4	4	Talbot	not_known	Male	33	1961-10-03 00:00:00	not_known	
5	5	Sheila-kathryn	Calton	Female	56	1977-05-13 00:00:00	Senior Editor	
...
3996	3996	Rosalia	Halgarth	Female	8	1975-08-09 00:00:00	VP Product Management	
3997	3997	Blanch	Nisuis	Female	87	2001-07-13 00:00:00	Statistician II	
3998	3998	Sarene	Woolley	U	60	NaN	Assistant Manager	
3999	3999	Patrizius	not_known	Male	11	1973-10-24 00:00:00	not_known	
4000	4000	Kippy	Oldland	Male	76	1991-11-05 00:00:00	Software Engineer IV	

4000 rows × 12 columns

```
In [41]: cd.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 1 to 4000
Data columns (total 12 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   customer_id                           4000 non-null   object
 1   first_name                             4000 non-null   object
 2   last_name                              4000 non-null   object
 3   gender                                 4000 non-null   object
 4   past_3_years_bike_related_purchases  4000 non-null   object
 5   DOB                                    3913 non-null   object
 6   job_title                             4000 non-null   object
 7   job_industry_category                 4000 non-null   object
 8   wealth_segment                        4000 non-null   object
 9   deceased_indicator                    4000 non-null   object
10   owns_car                              4000 non-null   object
11   tenure                                4000 non-null   int64
dtypes: int64(1), object(11)
memory usage: 375.1+ KB
```

Recasting the data types

```
In [42]: cd['DOB']=pd.to_datetime(cd['DOB'])
```

```
In [43]: cd=cd.astype({'past_3_years_bike_related_purchases':'int64',
                      'tenure':'float64','customer_id':"int64"

                      })
```

```
In [44]: cd.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 1 to 4000
Data columns (total 12 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   customer_id                           4000 non-null   int64
 1   first_name                             4000 non-null   object
 2   last_name                              4000 non-null   object
 3   gender                                 4000 non-null   object
 4   past_3_years_bike_related_purchases  4000 non-null   int64
 5   DOB                                    3913 non-null   datetime64[ns]
 6   job_title                             4000 non-null   object
 7   job_industry_category                 4000 non-null   object
 8   wealth_segment                        4000 non-null   object
 9   deceased_indicator                    4000 non-null   object
10   owns_car                              4000 non-null   object
11   tenure                                4000 non-null   float64
dtypes: datetime64[ns](1), float64(1), int64(2), object(8)
memory usage: 375.1+ KB
```

```
In [45]: cd.isnull().sum()
```

```
Out[45]: 0
customer_id
```

0

```

first_name      0
last_name       0
gender          0
past_3_years_bike_related_purchases  0
DOB            87
job_title       0
job_industry_category  0
wealth_segment  0
deceased_indicator  0
owns_car        0
tenure          0
dtype: int64

```

```
In [46]: cd.dropna(inplace=True)
```

Exploring the columns

```
In [47]: cd.columns
```

```
Out[47]: Index(['customer_id', 'first_name', 'last_name', 'gender',
               'past_3_years_bike_related_purchases', 'DOB', 'job_title',
               'job_industry_category', 'wealth_segment', 'deceased_indicator',
               'owns_car', 'tenure'],
              dtype='object', name=0)
```

```
In [48]: cd ['gender'].value_counts()
```

```
Out[48]: Female      2037
Male        1872
F           1
U           1
Femal       1
M           1
Name: gender, dtype: int64
```

```
In [49]: cd["gender"]=cd["gender"].str.replace("Femaleemale","Female")
```

```
In [50]: cd ['gender'].value_counts()
```

```
Out[50]: Female      2037
Male        1872
F           1
U           1
Femal       1
M           1
Name: gender, dtype: int64
```

```
In [51]: cd.drop(cd.index[(cd["gender"] == "F")],axis=0,inplace=True)
```

```
In [52]: cd.drop(cd.index[(cd["gender"] == "Femal")],axis=0,inplace=True)
```

```
In [53]: cd.drop(cd.index[(cd["gender"] == "M")],axis=0,inplace=True)
```

```
In [54]: cd ['gender'].value_counts()
```

```
Out[54]: Female      2037
Male        1872
U           1
Name: gender, dtype: int64
```

```
In [55]: cd["gender"]=cd["gender"].str.replace("U","others")
```

```

In [56]: cd ['gender'].value_counts()

Out[56]: Female      2037
         Male        1872
         others         1
         Name: gender, dtype: int64

In [57]: cd ['job_title'].value_counts()

Out[57]: not_known      497
         Tax Accountant   43
         Business Systems Development Analyst  43
         Social Worker    42
         Recruiting Manager  41
         ...
         Database Administrator I  4
         Health Coach I  3
         Health Coach III  3
         Research Assistant III  3
         Developer I  1
         Name: job_title, Length: 196, dtype: int64

In [58]: cd ['job_industry_category'].value_counts()

Out[58]: Manufacturing      796
         Financial Services  767
         not_known          655
         Health             595
         Retail             358
         Property           266
         IT                 152
         Entertainment      136
         Argiculture        113
         Telecommunications   72
         Name: job_industry_category, dtype: int64

In [59]: cd ['wealth_segment'].value_counts()

Out[59]: Mass Customer      1951
         High Net Worth     996
         Affluent Customer  963
         Name: wealth_segment, dtype: int64

In [60]: cd ['deceased_indicator'].value_counts()

Out[60]: N      3908
         Y        2
         Name: deceased_indicator, dtype: int64

In [61]: cd ['owns_car'].value_counts()

Out[61]: Yes      1971
         No       1939
         Name: owns_car, dtype: int64

In [62]: cd ['tenure'].value_counts()

Out[62]: 7.0      235
         5.0      228
         11.0     220
         10.0     218
         16.0     215
         8.0      211
         18.0     207
         12.0     202
         9.0      200
         14.0     200

```

```
6.0    192
4.0    191
13.0   190
17.0   182
15.0   179
1.0    166
3.0    160
19.0   159
2.0    150
20.0    96
22.0    55
21.0    54
Name: tenure, dtype: int64
```

Adding the Age Column

```
In [63]: from datetime import datetime, date as dt
```

```
cd["year"]=cd["DOB"].dt.year
today=dt.today()
cd.astype({"year":"int64"})
cd["age"]=today.year-cd["year"]
cd.drop("year",axis=1,inplace=True)
```

```
In [64]: cd
```

```
Out[64]:
```

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	job_id
2	2	Eli	Bockman	Male	81	1980-12-16	Administrative Officer	
3	3	Arlin	Dearle	Male	61	1954-01-20	Recruiting Manager	
4	4	Talbot	not_known	Male	33	1961-10-03	not_known	
5	5	Sheila-kathryn	Calton	Female	56	1977-05-13	Senior Editor	
6	6	Curr	Duckhouse	Male	35	1966-09-16	not_known	
...
3995	3995	Rusty	lapico	Male	93	1975-12-12	Staff Scientist	
3996	3996	Rosalia	Halgarth	Female	8	1975-08-09	VP Product Management	
3997	3997	Blanch	Nisuis	Female	87	2001-07-13	Statistician II	
3999	3999	Patrizius	not_known	Male	11	1973-10-24	not_known	
4000	4000	Kippy	Oldland	Male	76	1991-11-05	Software Engineer IV	

3910 rows × 13 columns

Customer address data

```
In [65]: ca=pd.read_excel(data,'CustomerAddress')
```

```
In [66]: ca
```

Out[66]:

Note: The data and information in this document is reflective of a hypothetical situation and client. This document is to be used for KPMG Virtual Internship purposes only.

0

customer_id

address

postcode

state

country

property_valuation

1	1	060 Morning Avenue	2016	New South Wales	Australia	10
2	2	6 Meadow Vale Court	2153	New South Wales	Australia	10
3	4	0 Holy Cross Court	4211	QLD	Australia	9
4	5	17979 Del Mar Point	2448	New South Wales	Australia	4
...
3995	3999	1482 Hauk Trail	3064	VIC	Australia	3
3996	4000	57042 Village Green Point	4511	QLD	Australia	6
3997	4001	87 Crescent Oaks Alley	2756	NSW	Australia	10
3998	4002	8194 Lien Street	4032	QLD	Australia	7
3999	4003	320 Acker Drive	2251	NSW	Australia	7

4000 rows × 6 columns

```
In [67]: ca.columns=ca.iloc[0]
```

```
In [68]: ca.drop(0,axis=0,inplace=True)
```

```
In [69]: ca.head()
```

Out[69]:

	customer_id	address	postcode	state	country	property_valuation
1	1	060 Morning Avenue	2016	New South Wales	Australia	10
2	2	6 Meadow Vale Court	2153	New South Wales	Australia	10
3	4	0 Holy Cross Court	4211	QLD	Australia	9

4	5	17979 Del Mar Point	2448	New South Wales	Australia	4
5	6	9 Oakridge Court	3216	VIC	Australia	9

```
In [70]: ca.nunique()
```

```
Out[70]: 0
customer_id      3999
address          3996
postcode         873
state            5
country          1
property_valuation 12
dtype: int64
```

```
In [71]: ca.duplicated().sum()
```

```
Out[71]: 0
```

No duplicates present

```
In [72]: ca.shape
```

```
Out[72]: (3999, 6)
```

```
In [73]: ca.describe()
```

```
Out[73]:
```

	customer_id	address	postcode	state	country	property_valuation
count	3999	3999	3999	3999	3999	3999
unique	3999	3996	873	5	1	12
top	1	3 Mariners Cove Terrace	2170	NSW	Australia	9
freq	1	2	31	2054	3999	647

```
In [74]: ca.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3999 entries, 1 to 3999
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customer_id           3999 non-null    object
1   address               3999 non-null    object
2   postcode              3999 non-null    object
3   state                 3999 non-null    object
4   country               3999 non-null    object
5   property_valuation    3999 non-null    object
dtypes: object(6)
memory usage: 187.6+ KB
```

Missing values

```
In [75]: ca.isnull().sum()
```

```
Out[75]: 0
customer_id      0
```



```

address      0
postcode     0
state        0
country      0
property_valuation  0
dtype: int64

```

```
In [76]: ca.isna().sum()
```

```

Out[76]: 0
customer_id      0
address          0
postcode         0
state            0
country          0
property_valuation  0
dtype: int64

```

No missing or na values

Recasting the data types

```
In [77]: ca=ca.astype({"customer_id":"int64","postcode":"int64","property_valuation":"int64"})
```

```
In [78]: ca.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3999 entries, 1 to 3999
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   customer_id           3999 non-null   int64
1   address                3999 non-null   object
2   postcode               3999 non-null   int64
3   state                  3999 non-null   object
4   country                3999 non-null   object
5   property_valuation     3999 non-null   int64
dtypes: int64(3), object(3)
memory usage: 187.6+ KB

```

Exploring the columns

```
In [79]: ca.columns
```

```

Out[79]: Index(['customer_id', 'address', 'postcode', 'state', 'country',
               'property_valuation'],
              dtype='object', name=0)

```

```
In [80]: ca['address'].value_counts()
```

```

Out[80]: 3 Mariners Cove Terrace      2
3 Talisman Place                    2
64 Macpherson Junction              2
359 Briar Crest Road                1
4543 Service Terrace                1
..
5063 Shopko Pass                    1
09 Hagan Pass                       1
87897 Lighthouse Bay Pass           1

```

```
294 Lawn Junction          1
320 Acker Drive            1
Name: address, Length: 3996, dtype: int64
```

```
In [81]: ca['state'].value_counts()
```

```
Out[81]: NSW          2054
VIC            939
QLD           838
New South Wales  86
Victoria       82
Name: state, dtype: int64
```

Changing the names, Victoria to VIC and New South Wales to NSW

```
In [82]: ca['state'].replace('New South Wales', 'NSW', inplace=True)
ca['state'].replace('Victoria', 'VIC', inplace=True)
ca.dropna(inplace=True)
ca
```

```
Out[82]:
```

	customer_id	address	postcode	state	country	property_valuation
1	1	060 Morning Avenue	2016	NSW	Australia	10
2	2	6 Meadow Vale Court	2153	NSW	Australia	10
3	4	0 Holy Cross Court	4211	QLD	Australia	9
4	5	17979 Del Mar Point	2448	NSW	Australia	4
5	6	9 Oakridge Court	3216	VIC	Australia	9
...
3995	3999	1482 Hauk Trail	3064	VIC	Australia	3
3996	4000	57042 Village Green Point	4511	QLD	Australia	6
3997	4001	87 Crescent Oaks Alley	2756	NSW	Australia	10
3998	4002	8194 Lien Street	4032	QLD	Australia	7
3999	4003	320 Acker Drive	2251	NSW	Australia	7

3999 rows × 6 columns

```
In [83]: ca['state'].value_counts()
```

```
Out[83]: NSW      2140
VIC       1021
QLD       838
Name: state, dtype: int64
```

All the columns have correct information.

TASK -2

Sprocket Central Pty Ltd has given us a new list of 1000 potential customers with their demographics and

attributes.

```
In [84]: Ncl=pd.read_excel(data,"NewCustomerList")
```

```
In [85]: Ncl
```

Out[85]:

Note: The data and information in this document is reflective of a hypothetical situation and client. This document is to be used for KPMG Virtual Internship purposes only.

Unnamed: 1

Unnamed: 2

Unnamed: 3

Unnamed: 4

Unnamed: 5

0	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	job_in
1	Chickie	Brister	Male	86	1957-07-12	General Manager	
2	Morly	Genery	Male	69	1970-03-22	Structural Engineer	
3	Ardelis	Forrester	Female	10	1974-08-28 00:00:00	Senior Cost Accountant	Fi
4	Lucine	Stutt	Female	64	1979-01-28	Account Representative III	
...	
996	Ferdinand	Romanetti	Male	60	1959-10-07	Paralegal	Fi
997	Burk	Wortley	Male	22	2001-10-17	Senior Sales Associate	
998	Melloney	Temby	Female	17	1954-10-05	Budget/Accounting Analyst IV	Fi
999	Dickie	Cubbini	Male	30	1952-12-17	Financial Advisor	Fi
1000	Sylas	Duffill	Male	56	1955-10-02	Staff Accountant IV	

1001 rows × 23 columns

```
In [86]: Ncl.columns=Ncl.iloc[0]
```

```
In [87]: Ncl.drop(0,axis=0,inplace=True)
```

```
In [88]: Ncl.head()
```

Out[88]:

	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	job_industry_catego
1	Chickie	Brister	Male	86	1957-07-12	General Manager	Manufacturir
2	Morly	Genery	Male	69	1970-03-22	Structural Engineer	Proper
3	Ardelis	Forrester	Female	10	1974-08-28 00:00:00	Senior Cost Accountant	Financial Servic
4	Lucine	Stutt	Female	64	1979-01-28	Account Representative III	Manufacturir
5	Melinda	Hadlee	Female	34	1965-09-21	Financial Analyst	Financial Servic

5 rows × 23 columns

Dropping unknown columns

In [89]:

Ncl.drop(Ncl.columns[[16,17,18,19,20]], axis=1, inplace=True)

In [90]:

Ncl.shape

Out[90]: (1000, 18)

In [91]:

Ncl.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 1 to 1000
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   first_name                           1000 non-null   object
1   last_name                             971 non-null    object
2   gender                               1000 non-null   object
3   past_3_years_bike_related_purchases 1000 non-null   object
4   DOB                                   983 non-null    object
5   job_title                             894 non-null    object
6   job_industry_category                835 non-null    object
7   wealth_segment                       1000 non-null   object
8   deceased_indicator                   1000 non-null   object
9   owns_car                             1000 non-null   object
10  tenure                               1000 non-null   object
11  address                              1000 non-null   object
12  postcode                             1000 non-null   object
13  state                                1000 non-null   object
14  country                              1000 non-null   object
15  property_valuation                   1000 non-null   object
16  Rank                                 1000 non-null   object
17  Value                                1000 non-null   object
dtypes: object(18)
memory usage: 140.8+ KB
```

In [92]:

Ncl.describe().T

Out[92]:

	count	unique	top	freq
0				

first_name	1000	940	Rozamond	3
last_name	971	961	Sissel	2
gender	1000	3	Female	513
past_3_years_bike_related_purchases	1000	100	60	20
DOB	983	961	1965-07-03	2
job_title	894	184	Associate Professor	15
job_industry_category	835	9	Financial Services	203
wealth_segment	1000	3	Mass Customer	508
deceased_indicator	1000	1	N	1000
owns_car	1000	2	No	507
tenure	1000	23	9	79
address	1000	1000	45 Shopko Center	1
postcode	1000	522	2145	9
state	1000	3	NSW	506
country	1000	1	Australia	1000
property_valuation	1000	16	9	173
Rank	1000	324	760	13
Value	1000.0	324.0	0.6375	13.0

In [93]: Ncl.nunique()

Out[93]:

0	
first_name	940
last_name	961
gender	3
past_3_years_bike_related_purchases	100
DOB	961
job_title	184
job_industry_category	9
wealth_segment	3
deceased_indicator	1
owns_car	2
tenure	23
address	1000
postcode	522
state	3
country	1
property_valuation	16
Rank	324
Value	324
dtype: int64	

In [94]: Ncl.isnull().sum()

Out[94]:

0	
first_name	0
last_name	29
gender	0
past_3_years_bike_related_purchases	0
DOB	17
job_title	106
job_industry_category	165

```
wealth_segment      0
deceased_indicator  0
owns_car            0
tenure             0
address            0
postcode           0
state              0
country            0
property_valuation  0
Rank               0
Value              0
dtype: int64
```

```
In [95]: Ncl["last_name"].fillna("Not known", inplace=True)
Ncl["job_title"].fillna("Not known", inplace=True)
Ncl["job_industry_category"].fillna("Not known", inplace=True)
```

```
In [96]: Ncl.isnull().sum()
```

```
Out[96]:
```

0	
first_name	0
last_name	0
gender	0
past_3_years_bike_related_purchases	0
DOB	17
job_title	0
job_industry_category	0
wealth_segment	0
deceased_indicator	0
owns_car	0
tenure	0
address	0
postcode	0
state	0
country	0
property_valuation	0
Rank	0
Value	0
dtype: int64	

```
In [97]: Ncl['DOB'] = pd.to_datetime(Ncl["DOB"])
         Ncl.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 1 to 1000
Data columns (total 18 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   first_name                          1000 non-null   object
 1   last_name                           1000 non-null   object
 2   gender                              1000 non-null   object
 3   past_3_years_bike_related_purchases 1000 non-null   object
 4   DOB                                983 non-null    datetime64[ns]
 5   job_title                           1000 non-null   object
 6   job_industry_category               1000 non-null   object
 7   wealth_segment                      1000 non-null   object
 8   deceased_indicator                  1000 non-null   object
 9   owns_car                            1000 non-null   object
10   tenure                              1000 non-null   object
11   address                             1000 non-null   object
12   postcode                           1000 non-null   object
13   state                               1000 non-null   object
14   country                             1000 non-null   object
15   property_valuation                  1000 non-null   object
16   Rank                                1000 non-null   object
17   Value                               1000 non-null   object
```

```
dtypes: datetime64[ns](1), object(17)
memory usage: 140.8+ KB
```

```
In [98]: Ncl=Ncl.astype({"past_3_years_bike_related_purchases":"int64",
                        "tenure":"int64",
                        "postcode":"int64",
                        "property_valuation":"int64",
                        "Value":"float64" ,
                        "Rank":"int64",

                        })
```

```
In [99]: Ncl.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 1 to 1000
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   first_name                           1000 non-null   object
1   last_name                            1000 non-null   object
2   gender                               1000 non-null   object
3   past_3_years_bike_related_purchases 1000 non-null   int64
4   DOB                                  983 non-null    datetime64[ns]
5   job_title                            1000 non-null   object
6   job_industry_category                1000 non-null   object
7   wealth_segment                       1000 non-null   object
8   deceased_indicator                   1000 non-null   object
9   owns_car                             1000 non-null   object
10  tenure                               1000 non-null   int64
11  address                              1000 non-null   object
12  postcode                             1000 non-null   int64
13  state                                1000 non-null   object
14  country                              1000 non-null   object
15  property_valuation                   1000 non-null   int64
16  Rank                                 1000 non-null   int64
17  Value                                1000 non-null   float64
dtypes: datetime64[ns](1), float64(1), int64(5), object(11)
memory usage: 140.8+ KB
```

```
In [100... Ncl.duplicated().sum()
```

```
Out[100]: 0
```

No duplicated values found.

Exploring the columns

```
In [101... Ncl.columns
```

```
Out[101]: Index(['first_name', 'last_name', 'gender',
        'past_3_years_bike_related_purchases', 'DOB', 'job_title',
        'job_industry_category', 'wealth_segment', 'deceased_indicator',
        'owns_car', 'tenure', 'address', 'postcode', 'state', 'country',
        'property_valuation', 'Rank', 'Value'],
        dtype='object', name=0)
```

```
In [102... Ncl['gender'].value_counts()
```

```
Out[102]: Female    513
          Male      470
```

U 17
Name: gender, dtype: int64

```
In [103... Ncl["gender"]=Ncl["gender"].str.replace("U","others")
```

```
In [104... Ncl['gender'].value_counts()
```

```
Out[104]: Female      513  
Male        470  
others       17  
Name: gender, dtype: int64
```

```
In [105... Ncl.dropna(inplace=True)
```

```
In [106... from datetime import datetime,date as dt  
Ncl['Year'] = Ncl['DOB'].dt.year  
today=dt.today()  
today.year
```

```
Out[106]: 2023
```

```
In [107... Ncl.astype({"Year":"int64"})  
Ncl["age"]=today.year-Ncl["Year"]
```

```
In [108... Ncl
```

```
Out[108]:
```

	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	job_industry_c
1	Chickie	Brister	Male	86	1957-07-12	General Manager	Manuf
2	Morly	Genery	Male	69	1970-03-22	Structural Engineer	
3	Ardelis	Forrester	Female	10	1974-08-28	Senior Cost Accountant	Financial
4	Lucine	Stutt	Female	64	1979-01-28	Account Representative III	Manuf
5	Melinda	Hadlee	Female	34	1965-09-21	Financial Analyst	Financial
...	
996	Ferdinand	Romanetti	Male	60	1959-10-07	Paralegal	Financial
997	Burk	Wortley	Male	22	2001-10-17	Senior Sales Associate	
998	Melloney	Temby	Female	17	1954-10-05	Budget/Accounting Analyst IV	Financial
999	Dickie	Cubbini	Male	30	1952-12-17	Financial Advisor	Financial
1000	Sylas	Duffill	Male	56	1955-10-02	Staff Accountant IV	

Comparing the new customer and old customer table

```
In [109... import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(20,10))
plt.subplot(121)


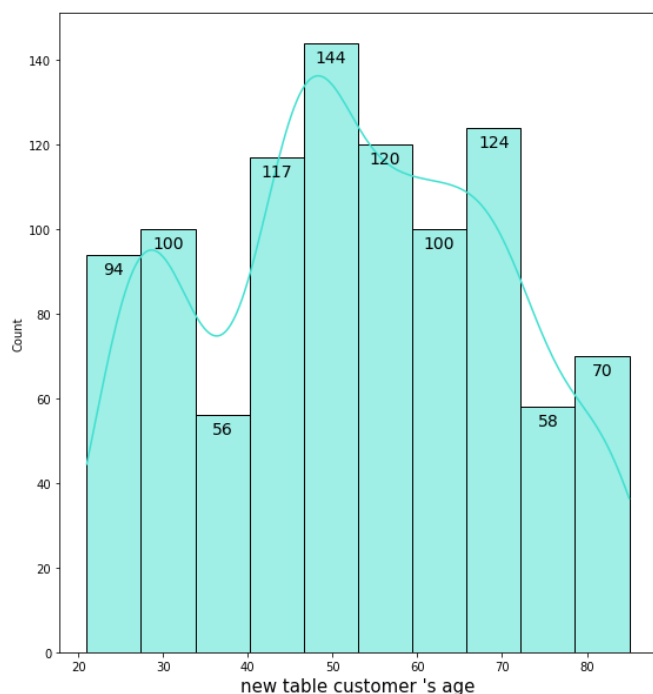
c=sns.histplot(data=Ncl,bins=10,x="age",color="turquoise",kde=True)
for i in c.patches:
    c.annotate(format(round(i.get_height()), '.0f'),
               (i.get_x() + i.get_width() / 2., i.get_height()),
               ha='center', va='center',color='black',
               size=14,
               xytext=(0, -12),
               textcoords='offset points')

plt.xlabel("new table customer 's age",fontsize=15)

plt.subplot(122)
d=sns.histplot(data=cd,bins=20,x="age",color="red",kde=True)
for i in d.patches:
    d.annotate(format(round(i.get_height()), '.0f'),
               (i.get_x() + i.get_width() / 2., i.get_height()),
               ha='center', va='center',color='black',
               size=14,
               xytext=(0, -12),
               textcoords='offset points')

plt.xlim(0, 100)
plt.xlabel(" old table customer's age",fontsize=15)
plt. suptitle("New vs Old customer 'age distribution",fontsize=20)
plt.show()
```

New vs Old customer 'age distribution



1)Most customer's age is in between 40-49 in new. In old, most of the people's age is also between 40-49. 2)The lowest group of age in old table is 80-100. 3)The lowest group in "new " table is 35-40. 4)There is a steep drop of customers in new table between 32-39. 5) Age group from 50-60 are considered most populated in both the tables.

Create a pivot table to make further visualisations

```
In [110]: table=pd.pivot_table(data=cd,index="gender",aggfunc="sum").reset_index("gender")
table
```

```
Out[110]:
```


	gender	age	customer_id	past_3_years_bike_related_purchases	tenure
0	Female	93693	4130807	98264	21708.0
1	Male	85843	3692409	93396	19931.0
2	others	180	34	59	20.0

```
In [111]: plt.figure(figsize=(10,5))
plt.subplot(121)
d=sns.barplot(data=table,x="gender",y="past_3_years_bike_related_purchases")
for i in d.patches:
    d.annotate(format(round(i.get_height()), '.0f'),
               (i.get_x() + i.get_width() / 2., i.get_height()),
               ha='center', va='center',color="black",
               size=10,
               xytext=(0, 5),
               textcoords='offset points')

plt.subplot(122)

plt.pie(data=table, x="past_3_years_bike_related_purchases", autopct='%0.2f%%', labels=["Fe
plt.suptitle("gender vs past three year purchases")
plt.title("pie chart")
```

```
Out[111]: Text(0.5, 1.0, 'pie chart')
```



Job distribution in old vs New table


```
In [112]: data=Ncl["job_industry_category"].value_counts()
```

data

```
Out[112]: Financial Services    202
Manufacturing    199
Not known        165
Health           152
Retail           78
Property         64
Entertainment    36
IT               36
Agriculture      26
Telecommunications 25
Name: job_industry_category, dtype: int64
```

```
In [113... keys=[202,199,165,152,78,64,36,36,26,25]
plt.figure(figsize=(10,10))
plt.pie(keys,labels=["Financial Services","Manufacturing", "Not known" , "Health " , "Re

plt.title("job category in New customer list")
plt.tight_layout()
```



1)20.55% of new customers are involved in Financial services which is the highest in new table. 2)Agriculture and Telecommunications are lowest in new customer table. 17% of jobs are still unidentified.

```
In [114... data=cd["job_industry_category"].value_counts()
data
```

```

Out[114]: Manufacturing      796
Financial Services      767
not_known              655
Health                 595
Retail                 358
Property               266
IT                     152
Entertainment          136
Agriculture            113
Telecommunications      72
Name: job_industry_category, dtype: int64


```

```

In [115... keys=[796,767,655,595,358,266,152,136,113,72]
labels=['Manufacturing ', 'Financial Services ', 'not_known ',
        'Health ', 'Retail', 'Property', 'IT', 'Entertainment ', 'Agriculture', 'Telecommunica
plt.figure(figsize=(10,10))
plt.pie(keys,labels=labels ,autopct="%.2f%%")

plt.title("job category in Old customer list")
plt.tight_layout()


```



Similar pattern obtained as of new list, here, the manufacturing consist of 20.36% and Financial services consist of 19.62%.

```
In [116... data=Ncl.groupby("wealth_segment").sum().reset_index()
```


```
In [117... plt.figure(figsize=(10,10))
sns.barplot(data=data,x='wealth_segment',y='past_3_years_bike_related_purchases')
plt.title('wealth_segment vs past 3 year purchases (new)')
plt.tight_layout()
```



Mass customer has the highest purchases in last three years. Affluent and high net worth customers some what share same records.

```
In [118... data=cd.groupby("wealth_segment").sum().reset_index()
```


```
In [119... plt.figure(figsize=(10,10))
sns.barplot(data=data,x='wealth_segment',y='past_3_years_bike_related_purchases')
plt.title('wealth_segment vs past 3 year purchases (old) ')
plt.tight_layout()
```



Affluent customers and High net worht customers have similar purchases in last three years in the old data.

```
In [120... plt.figure(figsize=(20,15))
plt.subplot(121)

c=sns.histplot(data=Ncl,bins=10,x="age",hue="wealth_segment",palette=['lightyellow','gre
for i in c.patches:
    c.annotate(format(round(i.get_height()), '.0f'),
                (i.get_x() + i.get_width() / 2., i.get_height()),
                ha='center', va='center',color='black',
                size=10,
                xytext=(0, -12),
                textcoords='offset points')
plt.title(' NEW CUSTOMER WEALTH DISTRIBUTION BY AGE')
plt.show()
```




```
In [121... plt.figure(figsize=(20,20))
plt.subplot(121)
np.sqrt(cd["age"])
c=sns.histplot(data=cd,bins=10,x="age",hue="wealth_segment",palette=['lightyellow','gree
for i in c.patches:
    c.annotate(format(round(i.get_height()), '.0f'),
                (i.get_x() + i.get_width() / 2., i.get_height()),
                ha='center', va='center',color='black',
                size=10,
```

```

        xytext=(0, -12),textcoords='offset points'
    )
plt.xlim(0,100)
plt.title("OLD CUSTOMER WEALTH DISTRIBUTION BY AGE")
plt.show()

```





In both the table, mass customers have high purchases and there age is ranging at 45-55.

```
In [122...] data=pd.DataFrame(Ncl.groupby('state')['owns_car'].value_counts())
data.rename(columns={"owns_car":"count"},inplace=True)
data.reset_index("state",inplace=True)

data.reset_index("owns_car",inplace=True)
data
```

Out[122]:


	owns_car	state	count
0	No	NSW	267
1	Yes	NSW	232
2	Yes	QLD	125
3	No	QLD	101
4	No	VIC	129
5	Yes	VIC	129

```
In [123...] range(6)
```

Out[123]: range(0, 6)


```
In [124...] plt.figure(figsize=(10,10))
c=sns.barplot(data=data,x='state',y='count',hue="owns_car")
plt.title("car owners in different states in new list",fontsize=15)
for i in c.patches:
    c.annotate(format(round(i.get_height()), '.0f'),
               (i.get_x() + i.get_width() / 2., i.get_height()),
               ha='center', va='center',color='black',
               size=10,
               xytext=(0, -12),textcoords='offset points'
               )

plt.tight_layout()
```



NSW has largest amount o people that do not owns a car. Victoria has these numbers spread evenly. QLD has relatively a high number of customers that owns a car.

RFM Analysis



Recency- How recent was the customer's last purchase? Customers who recently made a purchase will still have the product on their mind and are more likely to purchase or use the product again. Businesses often measure recency in days. But,

depending on the product, they may measure it in years, weeks or even hours.Frequency- How often did this customer make a purchase in a given period? Customers who purchased once are often are more likely to purchase again. Additionally, first time customers may be good targets for follow-up advertising to convert them into more frequent customers.Monetary- How much money did the customer spend in a given period? Customers who spend a lot of money are more likely to spend money in the future and have a high value to a business.

Creating a new column to know last purchase

```
In [125]: most_recent_purchase = Transactions['transaction_date'].max()  
Transactions['last_purchase_days_ago'] = most_recent_purchase - Transactions['transaction_date']  
Transactions['last_purchase_days_ago'] /= np.timedelta64(1, 'D')  
  
Transactions.head(25)
```

Out[125]:	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	product_line
1	1	2	2950	2017-02-25	False	Approved	Solex	Standard
2	2	3	3120	2017-05-21	True	Approved	Trek Bicycles	Standard
3	3	37	402	2017-10-16	False	Approved	OHM Cycles	Standard
4	4	88	3135	2017-08-31	False	Approved	Norco Bicycles	Standard
5	5	78	787	2017-10-01	True	Approved	Giant Bicycles	Standard
6	6	25	2339	2017-03-08	True	Approved	Giant Bicycles	Road
7	7	22	1542	2017-04-21	True	Approved	WeareA2B	Standard
8	8	15	2459	2017-07-15	False	Approved	WeareA2B	Standard
9	9	67	1305	2017-08-10	False	Approved	Solex	Standard
10	10	12	3262	2017-08-30	True	Approved	WeareA2B	Standard
11	11	5	1986	2017-01-17	False	Approved	Trek Bicycles	Mountain
12	12	61	2783	2017-01-05	True	Approved	OHM Cycles	Standard
13	13	35	1243	2017-02-26	True	Approved	Trek Bicycles	Standard
14	14	16	2717	2017-09-10	False	Approved	Norco Bicycles	Standard
15	15	12	247	2017-06-11	False	Approved	Giant Bicycles	Standard
16	16	3	2961	2017-10-10	False	Approved	Trek Bicycles	Standard
17	17	79	2426	2017-04-03	False	Approved	Norco Bicycles	Standard
18	18	33	1842	2017-06-02	False	Approved	Giant Bicycles	Standard
19	19	54	2268	2017-04-06	True	Approved	WeareA2B	Standard
20	20	25	3002	2017-01-28	True	Approved	Giant	Road

							Bicycles	
21	21	27	1582	2017-10-09	False	Approved	Trek Bicycles	Standard
22	22	37	595	2017-06-29	True	Approved	OHM Cycles	Standard
23	23	37	2001	2017-04-08	True	Approved	OHM Cycles	Standard
24	24	82	515	2017-10-18	False	Approved	Giant Bicycles	Road
25	25	89	2822	2017-06-11	False	Approved	WeareA2B	Touring

```
In [126... rfmTable = Transactions.groupby('customer_id').agg({
    'last_purchase_days_ago': lambda x: x.min(),
    'customer_id': lambda x: len(x),
    'profit': lambda x: x.sum()
})

rfmTable.rename(columns={
    'last_purchase_days_ago': 'recency',
    'customer_id': 'frequency',
    'profit': 'monetary_value'
}, inplace=True
)
```

```
In [127... rfmTable.shape
```

```
Out[127]: (3494, 3)
```

```
In [128... rfmTable.head()
```

```
Out[128]:
```

	recency	frequency	monetary_value
customer_id			
1	7.0	11	3016
2	128.0	3	2226
3	102.0	8	3363
4	195.0	2	221
5	16.0	6	2394

```
In [129... quartiles = rfmTable.quantile(q=[0.25,0.50,0.75])
quartiles
```

```
Out[129]:
```

	recency	frequency	monetary_value
0.25	17.0	4.0	1874.00
0.50	44.0	6.0	2891.50
0.75	85.0	7.0	4240.75

Giving the rfm score

```

In [130]: rfmTable['R_rank'] = rfmTable['recency'].rank(ascending=False)
rfmTable['F_rank'] = rfmTable['frequency'].rank(ascending=True)
rfmTable['M_rank'] = rfmTable['monetary_value'].rank(ascending=True)

rfmTable["R_rank_norm"]=(rfmTable['R_rank']/(rfmTable['R_rank'].max()))*100
rfmTable["F_rank_norm"]=(rfmTable['F_rank']/(rfmTable['F_rank'].max()))*100
rfmTable["M_rank_norm"]=(rfmTable['M_rank']/(rfmTable['M_rank'].max()))*100
rfmTable.drop(["R_rank", "F_rank", "M_rank"],axis=1,inplace=True)
rfmTable

```

Out[130]:

	recency	frequency	monetary_value	R_rank_norm	F_rank_norm	M_rank_norm
--	---------	-----------	----------------	-------------	-------------	-------------

customer_id

1	7.0	11	3016	89.988476	97.838534	52.575844
2	128.0	3	2226	12.921348	12.367592	33.714940
3	102.0	8	3363	19.014693	83.395362	59.645106
4	195.0	2	221	3.817344	4.308617	1.516886
5	16.0	6	2394	77.297609	57.171486	37.893532
...
3497	52.0	3	1649	44.281187	12.367592	20.105896
3498	127.0	6	3147	13.065399	57.171486	55.051517
3499	51.0	7	4957	44.929415	72.129974	84.559244
3500	144.0	6	1787	9.478536	57.171486	22.953635
5034	84.0	3	269	25.669836	12.367592	1.888952

3494 rows × 6 columns

RFM score is calculated based upon recency, frequency, monetary value normalize ranks. Based upon this score we divide our customers. Here we rate them on a scale of 5. Formula used for calculating rfm score is : $0.15 \times \text{Recency score} + 0.28 \times \text{Frequency score} + 0.57 \times \text{Monetary score}$

```

In [131]: rfmTable['RFM_Score'] = 0.15*rfmTable['R_rank_norm']+0.28 * \
          rfmTable['F_rank_norm']+0.57*rfmTable['M_rank_norm']
rfmTable['RFM_Score'] *= 0.05
rfmTable= rfmTable.round(2)
rfmTable.shape
rfmTable.reset_index("customer_id")
rfmTable.head(5)

```

Out[131]:

	recency	frequency	monetary_value	R_rank_norm	F_rank_norm	M_rank_norm	RFM_Score
--	---------	-----------	----------------	-------------	-------------	-------------	-----------

customer_id

1	7.0	11	3016	89.99	97.84	52.58	3.54
2	128.0	3	2226	12.92	12.37	33.71	1.23
3	102.0	8	3363	19.01	83.40	59.65	3.01
4	195.0	2	221	3.82	4.31	1.52	0.13
5	16.0	6	2394	77.30	57.17	37.89	2.46

Dividing the customers into categories :- Top Customer High value customer Medium Value Customer Low value customer Lost customer

```

In [132]: rfmTable["Customer_segment"] = np.where(rfmTable['RFM_Score'] >
          4.5, "Top Customers",

```

```

(np.where(
    rfmTable['RFM_Score'] > 4,
    "High value Customer",
    (np.where(
        rfmTable['RFM_Score'] > 3,
        "Medium Value Customer",
        np.where(rfmTable['RFM_Score'] > 1.6,
            'Low Value Customers', 'Lost Customers')))))
rfmTable[['RFM_Score', 'Customer_segment']].head(20)

```

Out[132]:


	RFM_Score	Customer_segment
customer_id		
1	3.54	Medium Value Customer
2	1.23	Lost Customers
3	3.01	Medium Value Customer
4	0.13	Lost Customers
5	2.46	Low Value Customers
6	2.86	Low Value Customers
7	0.23	Lost Customers
8	4.63	Top Customers
9	2.07	Low Value Customers
10	3.62	Medium Value Customer
11	3.03	Medium Value Customer
12	3.07	Medium Value Customer
13	3.68	Medium Value Customer
14	1.67	Low Value Customers
15	1.85	Low Value Customers
16	2.95	Low Value Customers
17	2.11	Low Value Customers
18	2.90	Low Value Customers
19	1.78	Low Value Customers
20	2.65	Low Value Customers

In [133...

```

plt.figure(figsize=(10,10))
sns.scatterplot(data=rfmTable,x=rfmTable['recency'],y=rfmTable['monetary_value'],hue="Cu
plt.title('Recency v/s Avg Profit')
plt.xlabel("Recency")
plt.ylabel('monetary value')
plt.show()

```



It shows that less recency customers haven't generated much of the monetary value. Lost customers haven't generated that much profits. Top customers have some High monetary value that even touches 12000 value but low recency. Medium value customers Coagulated at 3000-4000.

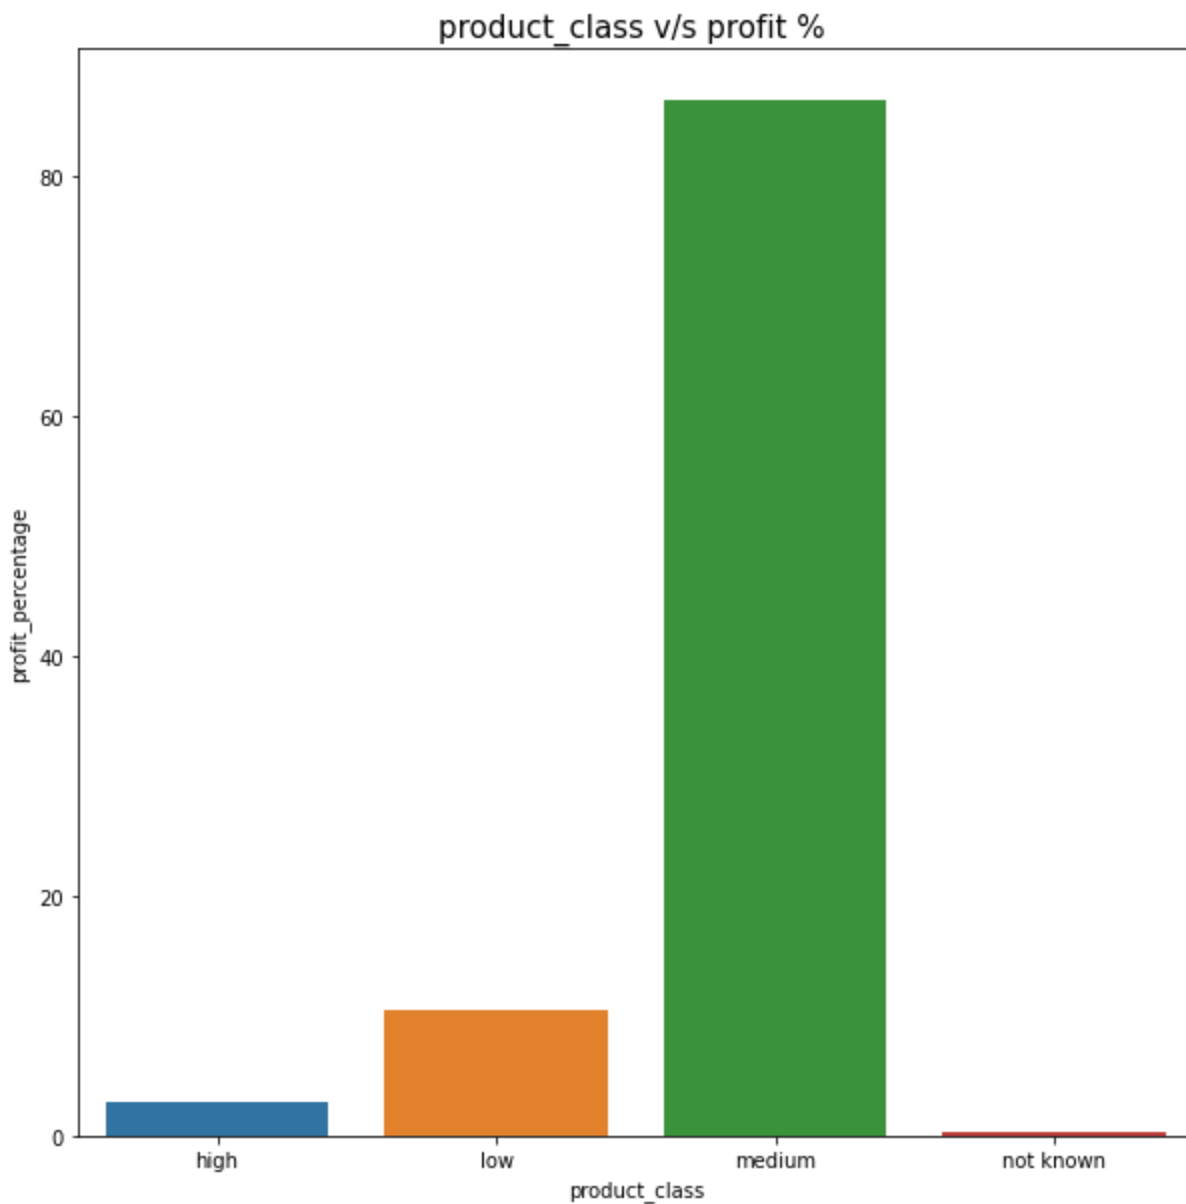
```
In [134]: data=Transactions.groupby("product_class").sum().reset_index("product_class")

data['profit_percentage']=(100*data['profit_percentage'])/data['profit_percentage'].sum(
data
```

Out[134]:	product_class	transaction_id	product_id	customer_id	list_price	standard_cost	profit	profit_percentage	la
0	high	30530037	120232	5236135	3287856	2171944	1115912	2.848345	
1	low	29355828	140918	5189023	2748093	1542882	1205211	10.551437	
2	medium	138024914	646143	24018100	15895665	7285718	8609947	86.303028	
3	not known	2099221	0	321663	214809	109532	105277	0.297190	

```
In [135]: plt.figure(figsize=(10,10))
c=sns.barplot(data=data,x='product_class',y="profit_percentage")
plt.title("product_class v/s profit % ",fontsize=15)
```

```
Out[135]: Text(0.5, 1.0, 'product_class v/s profit % ')
```



We can see that about 85 % of profit is been generated by the medium product_class. Somewhat 0.30% data is unknown. 10% profit is also obtained from the low product_class. Our target audience is basically who purchases the medium class product.

Merging the rfm Table and Old Customers Table

```
In [136... data=pd.merge(cd,rfmTable,on="customer_id",how="inner")
```

```
In [137... (data['Customer_segment']=="Top Customers").value_counts()
```

```
Out[137]: False    3273
          True     141
          Name: Customer_segment, dtype: int64
```

```
In [138... df=data.groupby("Customer_segment")["job_industry_category"].value_counts()
df1=pd.DataFrame(df)
df2=df1.iloc[40:]
df2.rename(columns={'job_industry_category':"count"},inplace=True)

df2.reset_index("job_industry_category",inplace=True)
df2
```

C:\Users\91913\AppData\Local\Temp\ipykernel_16880\1344750093.py:4: SettingWithCopyWarning:
g:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df2.rename(columns={'job_industry_category':"count"},inplace=True)
```

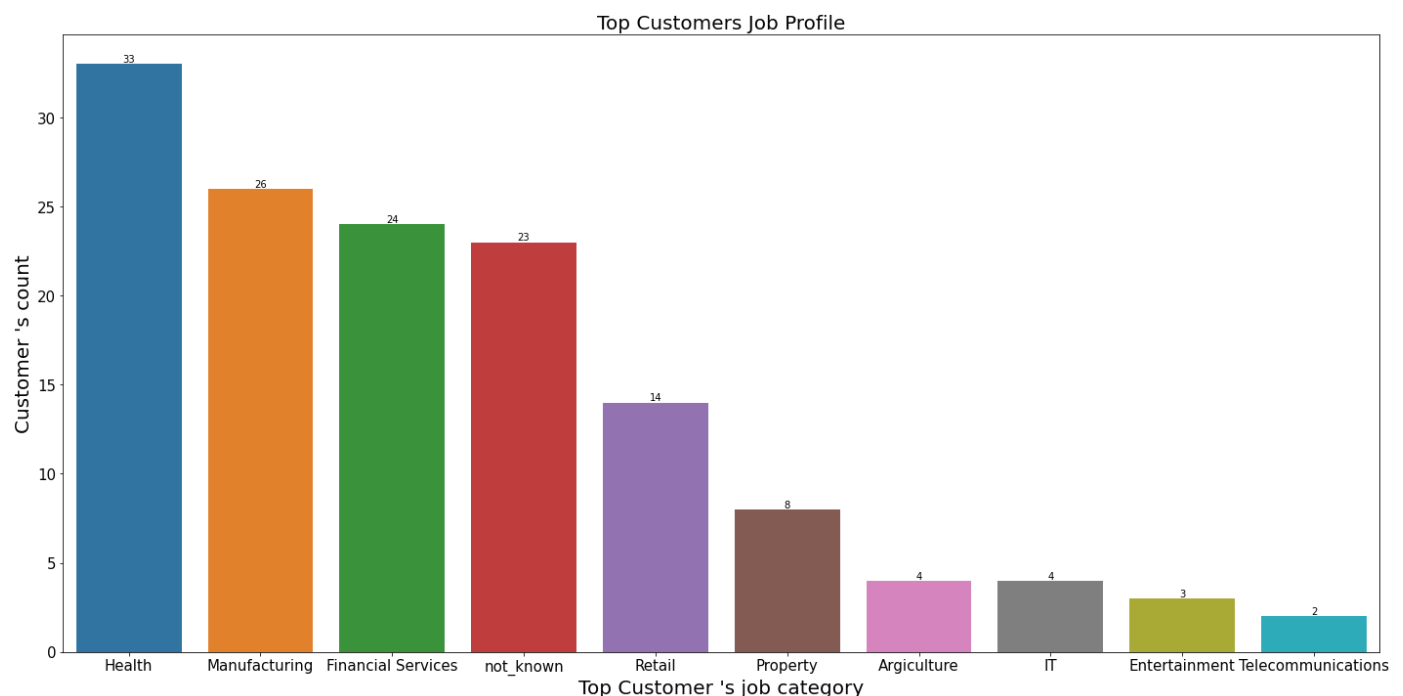
Out[138]:

job_industry_category		count
Customer_segment		
Top Customers	Health	33
Top Customers	Manufacturing	26
Top Customers	Financial Services	24
Top Customers	not_known	23
Top Customers	Retail	14
Top Customers	Property	8
Top Customers	Argiculture	4
Top Customers	IT	4
Top Customers	Entertainment	3
Top Customers	Telecommunications	2

In [139]...

```
plt.figure(figsize=(20,10))
d=sns.barplot(data=df2,x="job_industry_category",y="count")
plt.xlabel("Top Customer 's job category",fontsize=20)

for i in d.containers:
    d.bar_label(i,)
plt.ylabel("Customer 's count",fontsize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
plt.title("Top Customers Job Profile",fontsize=20)
plt.tight_layout()
```



Looking from the bar graph we get to know that the most of the top customer belong from the Health and Manufacturing sector. 23 customers also contribute but they are from unknown profiles.

In [140]...

```
df3=df1.iloc[10:20]
df3.rename(columns={'job_industry_category':"count"},inplace=True)
```

```
df3.reset_index("job_industry_category",inplace=True)
df3
```

```
C:\Users\91913\AppData\Local\Temp\ipykernel_16880\1467183473.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df3.rename(columns={'job_industry_category':"count"},inplace=True)
```

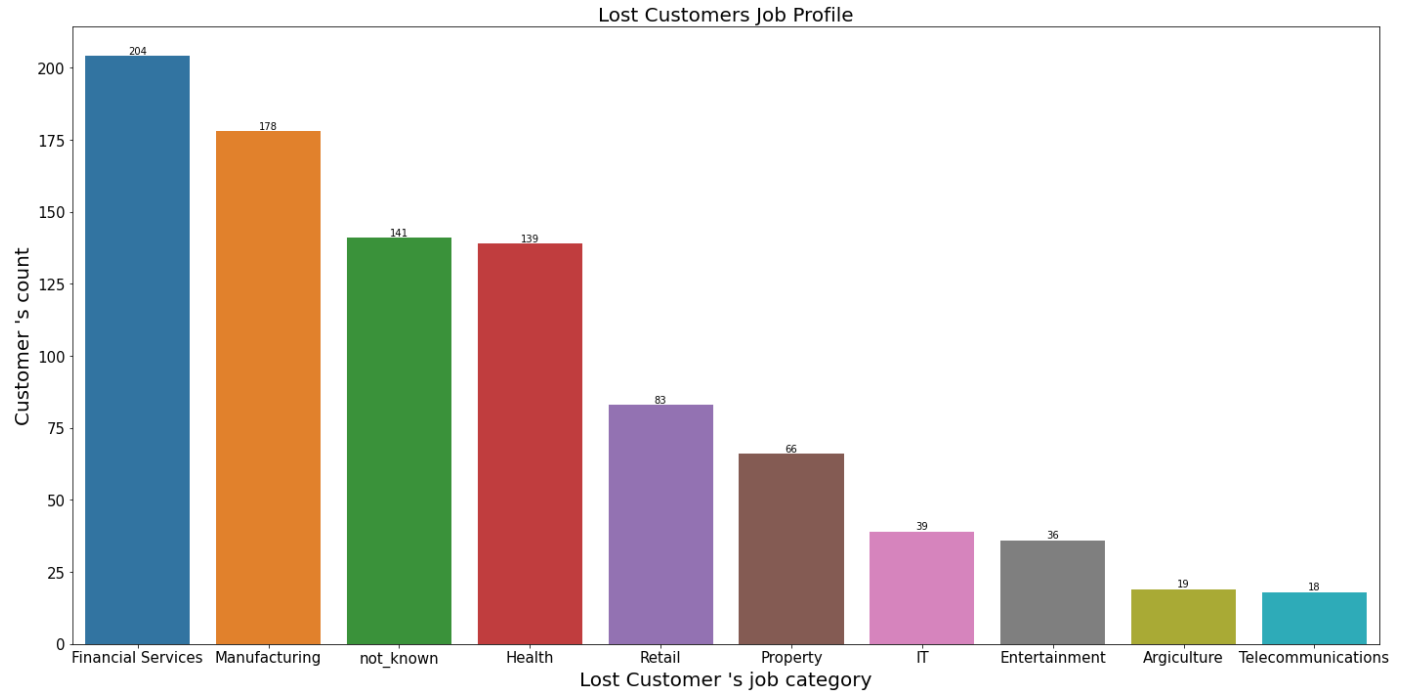
Out[140]:

job_industry_category		count
Customer_segment		
Lost Customers	Financial Services	204
Lost Customers	Manufacturing	178
Lost Customers	not_known	141
Lost Customers	Health	139
Lost Customers	Retail	83
Lost Customers	Property	66
Lost Customers	IT	39
Lost Customers	Entertainment	36
Lost Customers	Argiculture	19
Lost Customers	Telecommunications	18

In [141]...

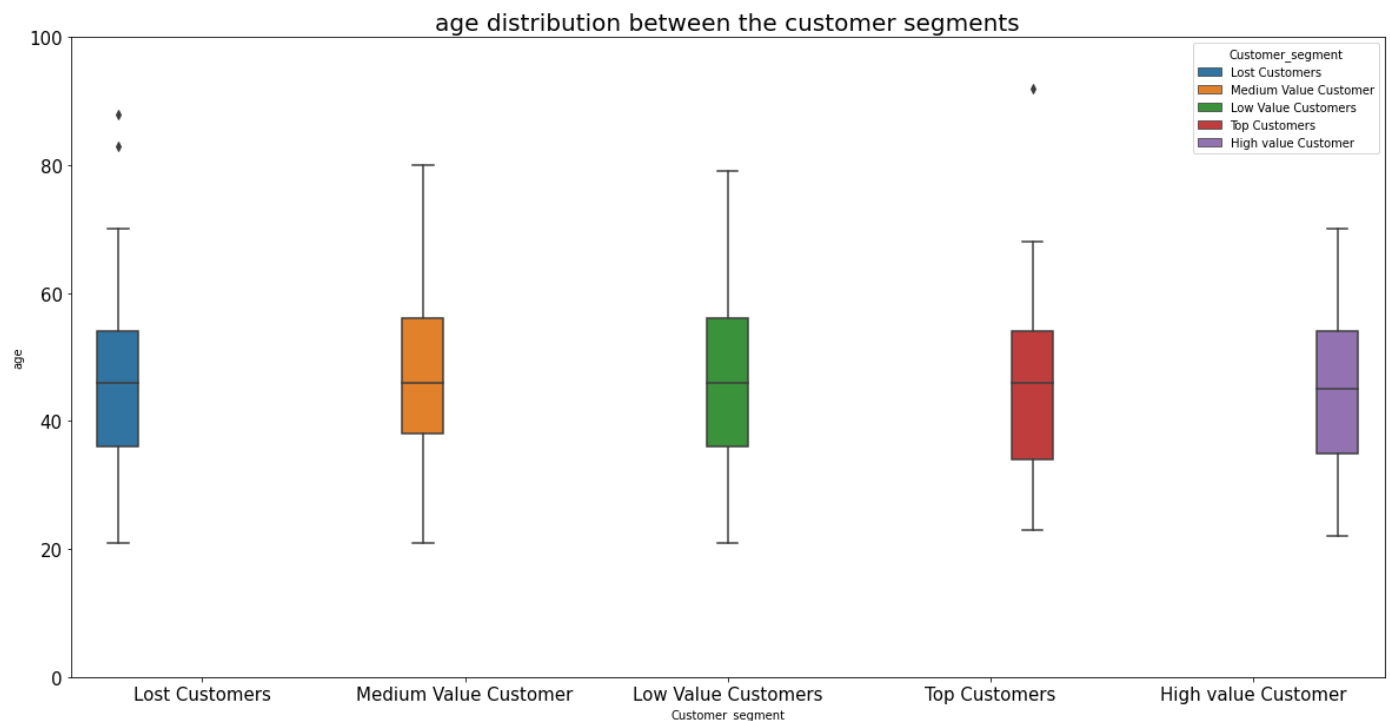
```
plt.figure(figsize=(20,10))
d=sns.barplot(data=df3,x="job_industry_category",y="count")
plt.xlabel("Lost Customer 's job category",fontsize=20)

for i in d.containers:
    d.bar_label(i,)
plt.ylabel("Customer 's count",fontsize=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
plt.title("Lost Customers Job Profile",fontsize=20)
plt.tight_layout()
```



The customers that we lost mostly belong from the Financial services background and manufacturing.

```
In [142... plt.figure(figsize=(20,10))
sns.boxplot(data=data,x="Customer_segment",y="age",hue="Customer_segment")
plt.xticks(fontsize=15)
plt.yticks(fontsize=15)
plt.ylim(0,100)
plt.title("age distribution between the customer segments",fontsize=20)
plt.show()
```



Most of the age of the customers lie between 40-50. Middle value customers have a high upper limit of approx 80 years.

Task -3

Tableau interactive Dashboard

1) Saving all the updated and corrected files to use as tableau data.

```
In [143... pip install openpyxl
```

Requirement already satisfied: openpyxl in c:\users\91913\anaconda3\lib\site-packages (3.0.9)

Requirement already satisfied: et-xmlfile in c:\users\91913\anaconda3\lib\site-packages (from openpyxl) (1.1.0)

Note: you may need to restart the kernel to use updated packages.

```
In [144... Transactions.to_excel("Transactions.xlsx")
Ncl.to_excel("Ncl.xlsx")
```

```
In [145... cd.to_excel("cd.xlsx")
```

```
In [146... ca.to_excel("ca.xlsx")
```

```
In [147... rfmTable.to_excel('rfm.xlsx')
%%html
<div class='tableauPlaceholder' id='viz1684846726063' style='position: relative'><noscri
```

